

# I/O SYSTEM CALLS

Alok Jhaldiyal

Assistant Professor

Dept. of Cyber Platform, SoCSE, UPES

# System Calls for I/O

- They look like regular procedure calls but are different
  - A system call makes a request to the operating system by trapping into kernel mode
  - A procedure call just jumps to a procedure defined elsewhere in your program
- Some library procedure calls may themselves make a system call
  - e.g., `fopen()` calls `open()`

# File Statistics

```
#include <sys/stat.h>
```

```
int stat(const char* name, struct stat* buf);
```

- Get information about a file
- Returns:
  - 0 on success
  - -1 on error, sets errno
- Parameters:
  - name: Path to file you want to use
    - Absolute paths begin with “/”, relative paths do not
  - buf: Statistics structure
    - off\_t st\_size: Size in bytes
    - time\_t st\_mtime: Date of last modification. Seconds since January 1, 1970

# File: Open

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open (const char* path, int flags [, int mode ]);
```


- Open (and/or create) a file for reading, writing or both
- Returns:
  - Return value  $\geq 0$  : Success - New file descriptor on success
  - Return value = -1: Error, check value of errno
- Parameters:
  - path: Path to file you want to use
    - Absolute paths begin with “/”, relative paths do not
  - flags: How you would like to use the file
  - O\_RDONLY: read only, O\_WRONLY: write only, O\_RDWR: read and write,
  - O\_CREAT: create file if it doesn't exist, O\_EXCL: prevent creation if it already exists

# Open: Example

```
#include <fcntl.h>
#include <errno.h>
extern int errno;
```

```
main() {
    int fd;
    fd = open("foo.txt", O_RDONLY | O_CREAT);
    printf("%d\n", fd);
    if (fd == -1) {
        printf("Error Number %d\n", errno);
        perror("Program");
    }
}
```

Argument: string  
Output: the string, a colon, and a description of the error condition stored in `errno`



# File: Close

```
#include <fcntl.h>
```

```
int close(int fd);
```

- Close a file
  - Tells the operating system you are done with a file descriptor
- Return:
  - 0 on success
  - -1 on error, sets errno
- Parameters:
  - fd: file descriptor

# Close: Example

```
#include <fcntl.h>
main() {
    int fd1;

    if(( fd1 = open("foo.txt", O_RDONLY)) < 0) {
        perror("c1");
        exit(1);
    }
    if (close(fd1) < 0) {
        perror("c1");
        exit(1);
    }
    printf("closed the fd.\n");
}
```

# File: Read

```
#include <fcntl.h>
```

```
size_t read (int fd, void* buf, size_t cnt);
```

- Read data from one buffer to file descriptor
  - Read **size** bytes from the file specified by **fd** into the memory location pointed to by **buf**
- Return: How many bytes were actually read
  - Number of bytes read on success
  - 0 on reaching end of file
  - -1 on error, sets **errno**
  - -1 on signal interrupt, sets **errno** to **EINTR**
- Parameters:
  - **fd**: file descriptor
  - **buf**: buffer to read data from
  - **cnt**: length of buffer



# File: Read

```
size_t read (int fd, void* buf, size_t cnt);
```

## ■ Things to be careful about

- **buf** needs to point to a valid memory location with length not smaller than the specified size
  - Otherwise, what could happen?
- **fd** should be a valid file descriptor returned from **open()** to perform read operation
  - Otherwise, what could happen?
- **cnt** is the requested number of bytes read, while the return value is the actual number of bytes read
  - How could this happen?

# Read: Example

```
#include <fcntl.h>
main() {
    char *c;
    int fd, sz;

    c = (char *) malloc(100
                        * sizeof(char));
    fd = open("foo.txt",
              O_RDONLY);
    if (fd < 0) {
        perror("r1");
        exit(1);
    }
```

```
sz = read(fd, c, 10);
printf("called
      read(%d, c, 10).
      returned that %d
      bytes were
      read.\n", fd, sz);
c[sz] = '\0';

printf("Those bytes
      are as follows:
      %s\n", c);
close(fd);
```

```
}
```

# File: Write

```
#include <fcntl.h>
```

```
size_t write (int fd, void* buf, size_t cnt);
```

- Write data from file descriptor into buffer
  - Writes the bytes stored in **buf** to the file specified by **fd**
- Return: How many bytes were actually written
  - Number of bytes written on success
  - 0 on reaching end of file
  - -1 on error, sets **errno**
  - -1 on signal interrupt, sets **errno** to **EINTR**
- Parameters:
  - **fd**: file descriptor
  - **buf**: buffer to write data to
  - **cnt**: length of buffer

# File: Write

```
size_t write (int fd, void* buf, size_t cnt);
```

- Things to be careful about
  - The file needs to be opened for write operations
  - **buf** needs to be at least as long as specified by **cnt**
    - If not, what will happen?
  - **cnt** is the requested number of bytes to write, while the return value is the actual number of bytes written
    - How could this happen?

# Write: Example

```
#include <fcntl.h>
```

```
main()
```

```
{
```

```
    int fd, sz;
```

```
    fd = open("out3",  
              O_RDWR | O_CREAT |  
              O_APPEND, 0644);
```

```
    if (fd < 0) {  
        perror("r1");  
        exit(1);  
    }
```

```
}
```

```
sz = write(fd, "cs241\n",  
           strlen("cs241\n"));
```

```
printf("called write(%d,  
       \"cs360\\n\", %d).  
       it returned %d\n",  
       fd, strlen("cs360\n"),  
       sz);
```

```
close(fd);
```

```
}
```

# Unix Error Model

## ■ Error Model

- **errno** variable
  - Unix provides a globally accessible integer variable that contains an error code number
- Return value
  - 0 on success
  - -1 on failure for functions returning integer values
  - NULL on failure for functions returning pointers
- Examples (see **errno.h**)

```
#define EPERM      1      /* Operation not permitted */
#define ENOENT     2      /* No such file or directory */
#define ESRCH     3      /* No such process */
#define EINTR     4      /* Interrupted system call */
#define EIO       5      /* I/O error */
#define ENXIO     6      /* No such device or address */
```