

Process Management

Alok Jhaldiyal
Assistant Professor
SoCSE, UPES
Dehradun

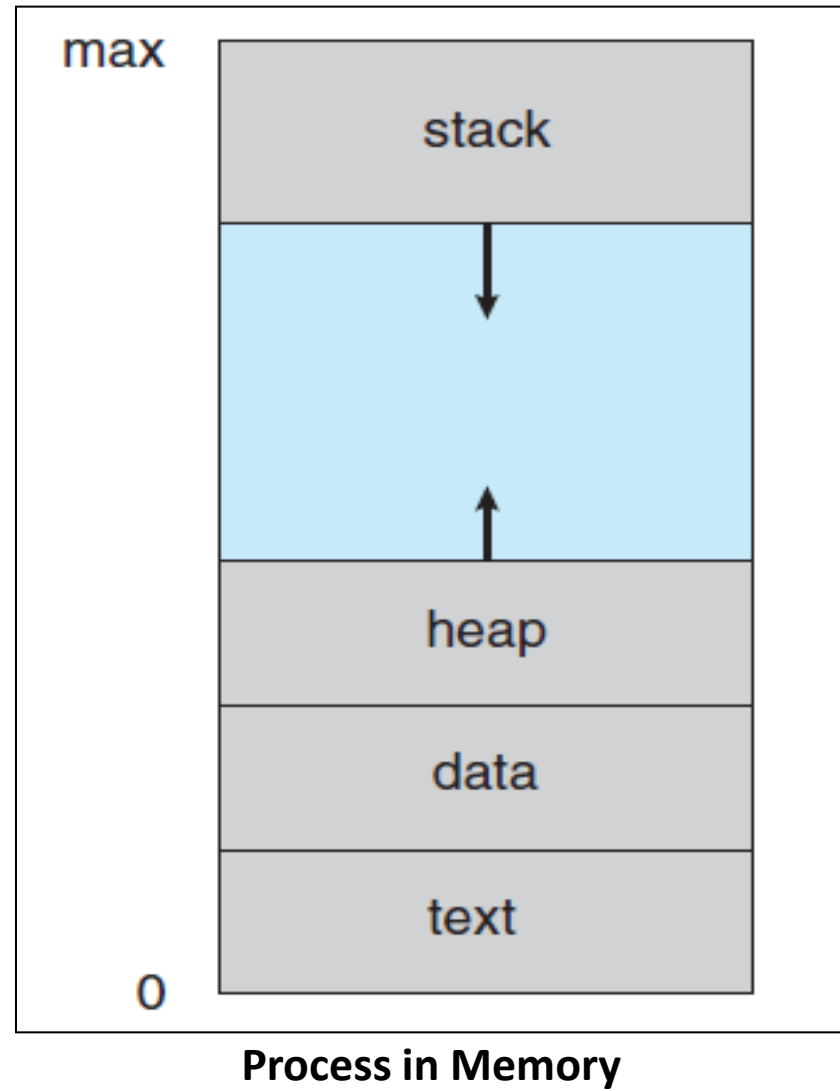
What a process is ?

- Process is
 - A program in execution
 - Unit of work of most operating systems
 - Requires resources – CPU time, memory, I/O devices
- A system may consist of multiple running processes: System process running system code and user process running user code.
- Traditionally process used to have a single thread but modern operating system contains multiple thread.
- Important aspects of process management: Thread management, creation and deletion of process, scheduling process and maintain a synchronized and a communicated state.

The Process

- A program is set of instruction stored in computers memory.
- A program may be active or in a passive state.
- Program in active state is a process.
- Process includes a program counter and data stored in the processor registers.
- A process generally includes process stack which contains temporary data, a data section, and also a heap.

- Process Stack contains temporary data such as function parameters, return addresses, and local variables.
- Data section contains global variables
- Heap is the dynamically allocated memory during runtime.
- And text is program instructions

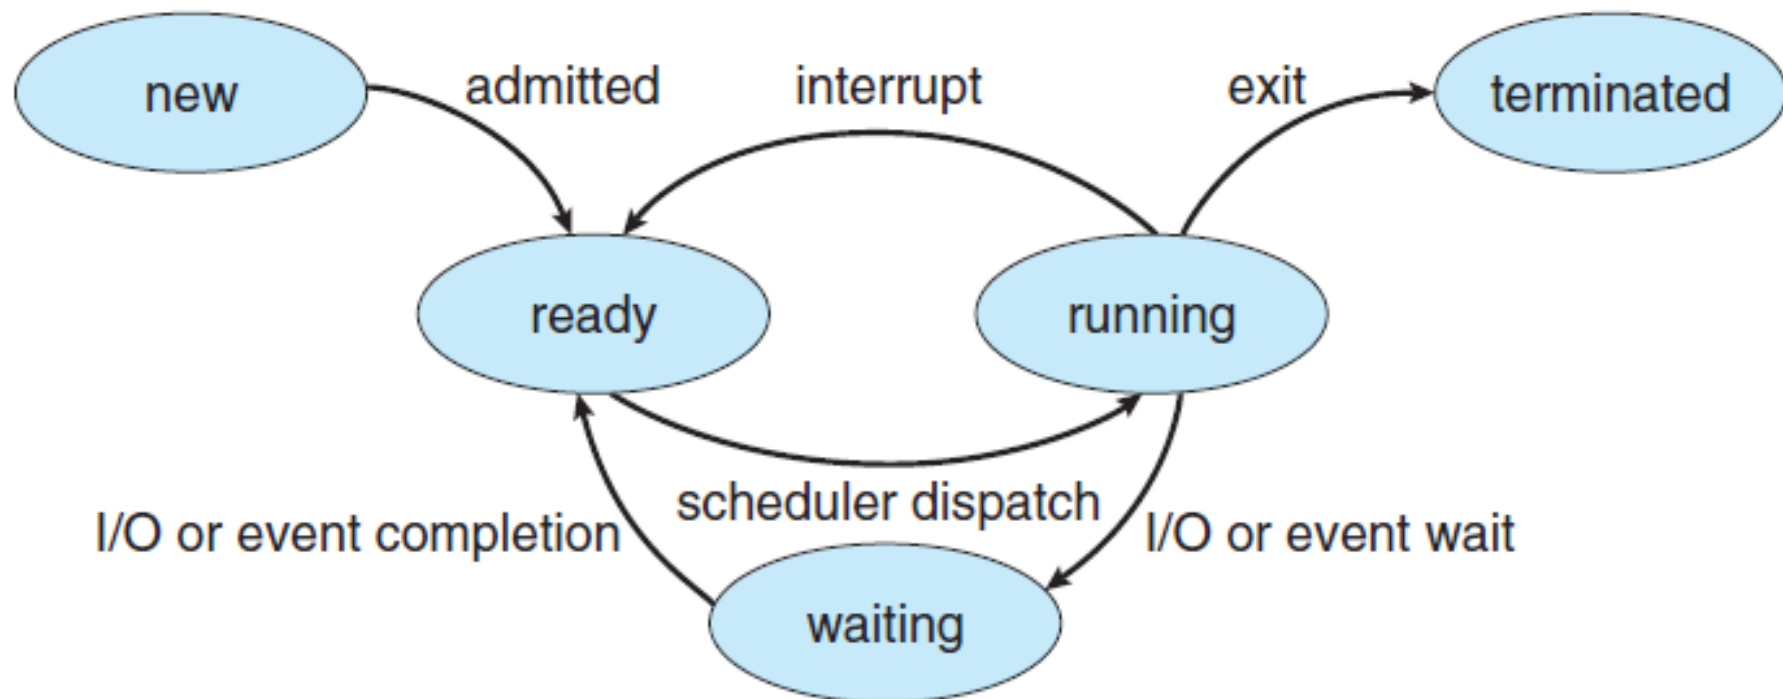


More about process

- A program becomes a process when an executable file is loaded into memory.
- Two processes may be associated with the same program, they are nevertheless considered two separate processes.
- Example can be a browser in which user has opened multiple tabs, but active tab is a separate process.

Process State

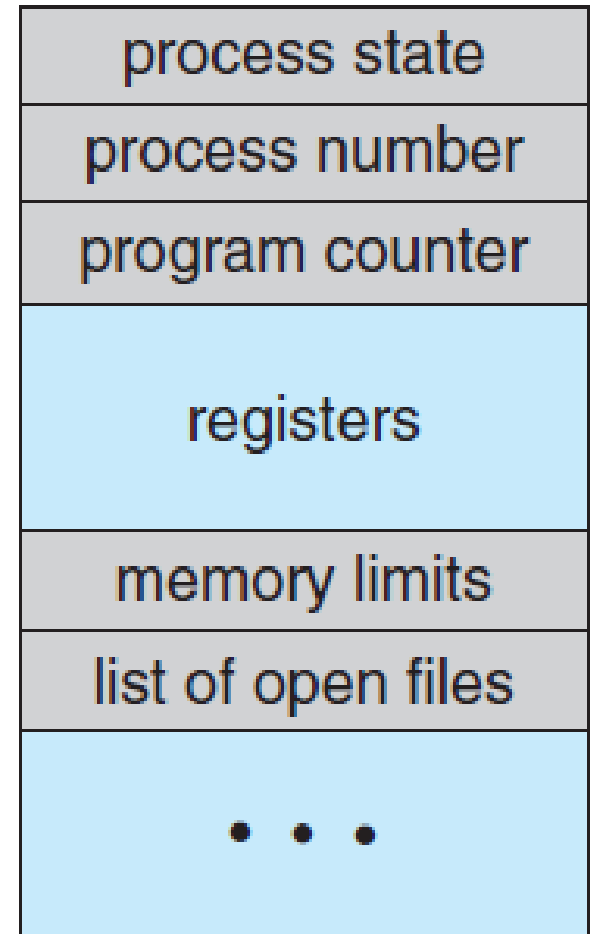
- When a process executes, it is in some state that we call process state.
- A process may be in one of the following states:
 - New. The process is being created.
 - Running. Instructions are being executed.
 - Waiting. The process is waiting for some event to occur (such as an I/O completion or reception of a signal).
 - Ready. The process is waiting to be assigned to a processor.
 - Terminated. The process has finished execution



Process State Diagram

Process Control Block

- Representation of a process in operating system is done using a Process Control Block (PCB) or Task Control Block.
- A PCB is a table and each block holding information associated with the particular process.

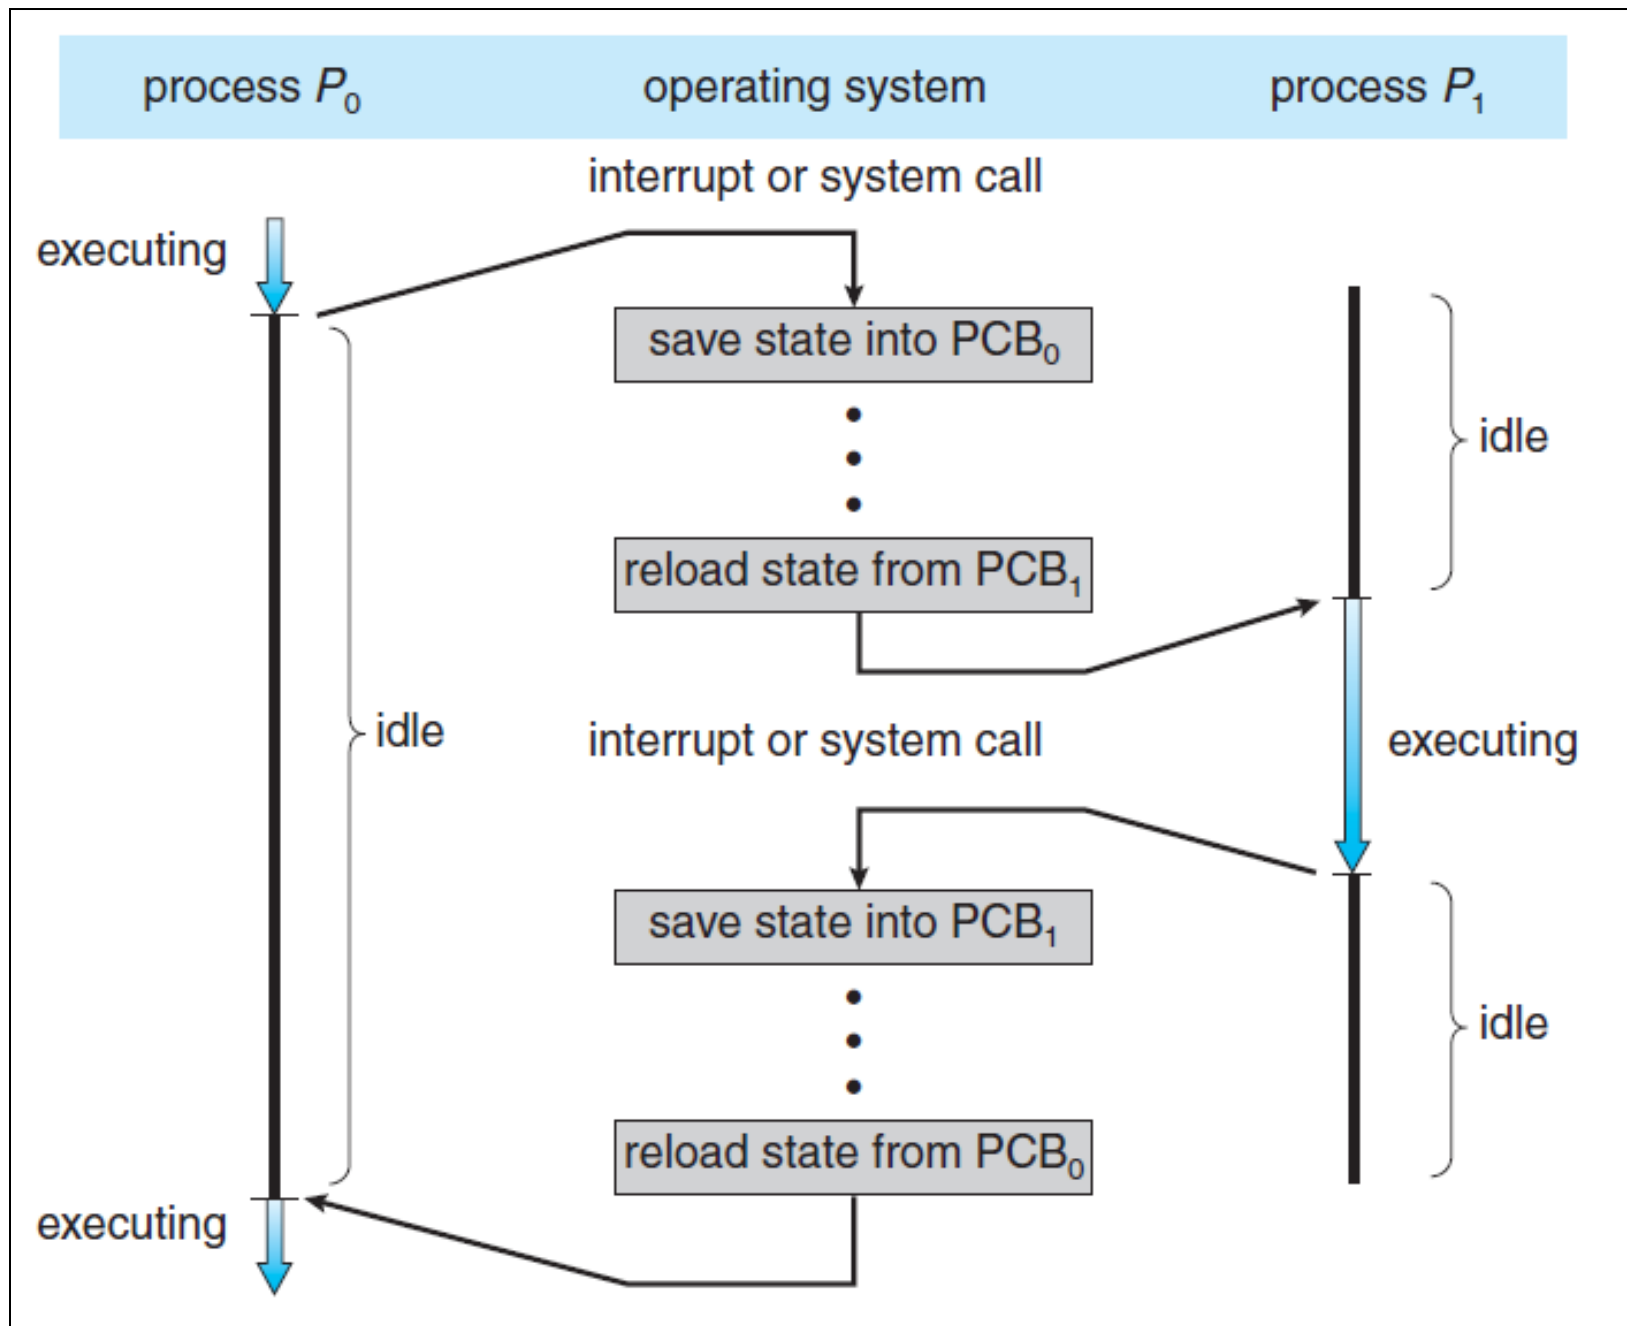


Process Control Block

Process Control Block

- **Pieces of information in PCB:**
 - **Process state:** New, Ready, Running, Waiting, Halted and Terminated.
 - **Program counter:** Indicates the address of the next instruction.
 - **CPU registers:**
 - Vary in type and number
 - Includes accumulators, index registers, stack pointers, and general-purpose registers
 - During interrupt process state and program counter information is stored in registers.
 - **CPU-scheduling information:** Includes a process priority, pointers to scheduling queues, and other scheduling parameters.

- **Memory-management information:** Might include value of the base and limit registers and the page tables, based on the memory system used by the underlying system.
- **Accounting information:** This information includes the amount of CPU and real time used, time limits, account numbers, job or process numbers.
- **I/O status information:** Includes the list of I/O devices allocated to the process and list of open files.



CPU Switching between two concurrent processes

Class Activity: Assume that a process is running in which user is recording his voice and at the same time it is getting converted into text form. Based on the above scenario:

- 1. List all the required resources in context to operating system to run this process successfully.**
- 2. Assume that due to system calls this process suffers two interrupts, using a diagram show the change in state on one side and highlight the changes in the PCB on the other side of your diagram.**

Context Switching

- Context Switch: A context switch is the mechanism to store and restore the state or context of a CPU in Process Control block (PCB) so that a process execution can be resumed from the same point at a later time.
- Process to switch the CPU between multiple programs which want to run on the system is context switching.
- Act of context switching can only be performed in kernel mode as only in kernel mode page map can be changed.

Context Switching

- Get into kernel mode.
- Save the old program's registers somewhere, and the address of the next instruction it was about to execute, so they can be restored later, save in its PCB.
- Save the mappings in the current page map, also into the PCB.
- Unmap all of the old program's pages from the page map.
- Choose a new program to re-start. Find its PCB.
- Re-map the pages of the new program from its PCB.
- Re-load the registers of the new program from its PCB.
- Return to the next instruction of the new program, and return to user-mode at the same time.

Process Scheduling

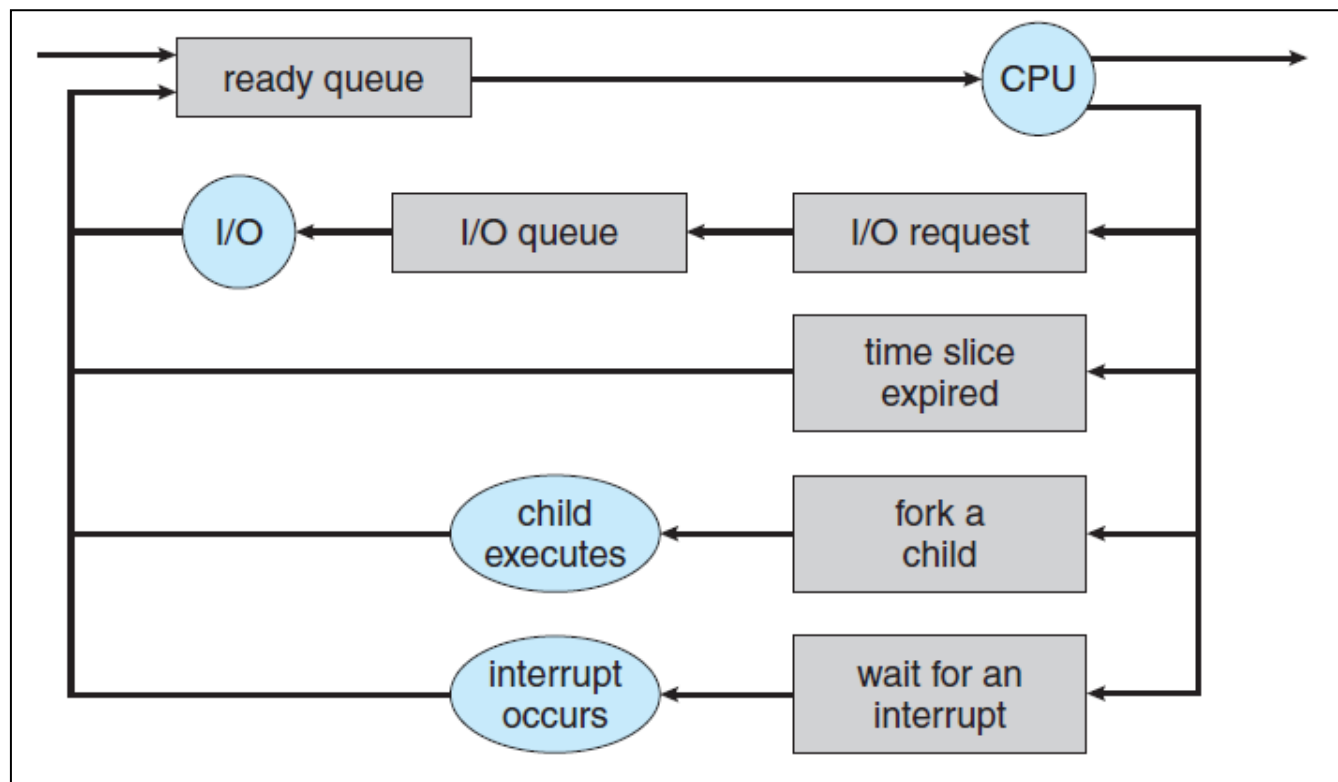
- Multiprogramming aims to maximize CPU utilization and it requires CPU to have some process running at all times, to maximize CPU utilization.
- Idea is to switch between CPU so frequently that the CPU can interact with each program.
- A process scheduler helps CPU find a available process.

Scheduling Queues

- All incoming jobs are kept in a job queue.
- Processes residing in main memory and are ready to be executed are kept in ready queue.
- Ready queue is a linked list and headers contain pointers to the first and last process PCB.
- Each PCB has a header field that holds the pointer to the next process PCB.
- A process waiting for a shared device is kept in a device queue.
- Each device has its own device queue and an oval denotes a resource.

Queuing Diagram

- A common representation of process scheduling is a queueing diagram.
- A queueing diagram is made of ovals and rectangles and connected using arrows.
- Two types of queues are present: the ready queue and a set of device queues.
- Rectangular box represents a queue



- A new process waits in a ready queue.
- Once process is assigned a CPU, it is executing.
- Once a process is executing several events may occur:
 - The process could issue an I/O request and then be placed in an I/O queue.
 - The process could create a new child process and wait for the child's termination.
 - The process could be removed forcibly from the CPU, as a result of an interrupt, and be put back in the ready queue.

Schedulers

- A process since its creation passes through several queues.
- A scheduler's responsibility is to select an appropriate and available process.
- Before getting a CPU time a process is spooled in memory disk.
- Scheduler types:
 - Long term scheduler
 - Short term scheduler
- The two schedulers are distinguished by the frequency of their operation.

Schedulers

- Short term scheduler:
 - It is CPU scheduler
 - Selects from among the processes that are ready to execute and allocates the CPU to one of them.
 - A CPU might run for few milliseconds before waiting for a input.
 - A short term scheduler needs to be very fast to select a process and assign CPU.
 - It often switches process at every 100 milliseconds.

Schedulers

- Long term schedulers:
 - These schedulers function less frequently, often minutes may separate two process.
 - Long term schedulers control the degree of multiprogramming in the system i.e. process loaded in the memory.
 - The rate to process creation must match the process departure rate.
 - Long term schedulers have flexibility to give more time to select a process.

Schedulers

- For efficiently utilizing CPU time a long term scheduler should select a proper mix of processes.
- A process is usually Input bound or CPU bound:
 - Input Bound: It is a process that spends more of its execution time waiting for inputs.
 - CPU Bound: A CPU bound generates I/O requests infrequently, using more of its time doing computations
- If all processes are I/O bound, the ready queue will almost always be empty, short term scheduler will have very less to do.
- If all processes are CPU bound, the I/O waiting queue will almost always be empty, devices will go unused

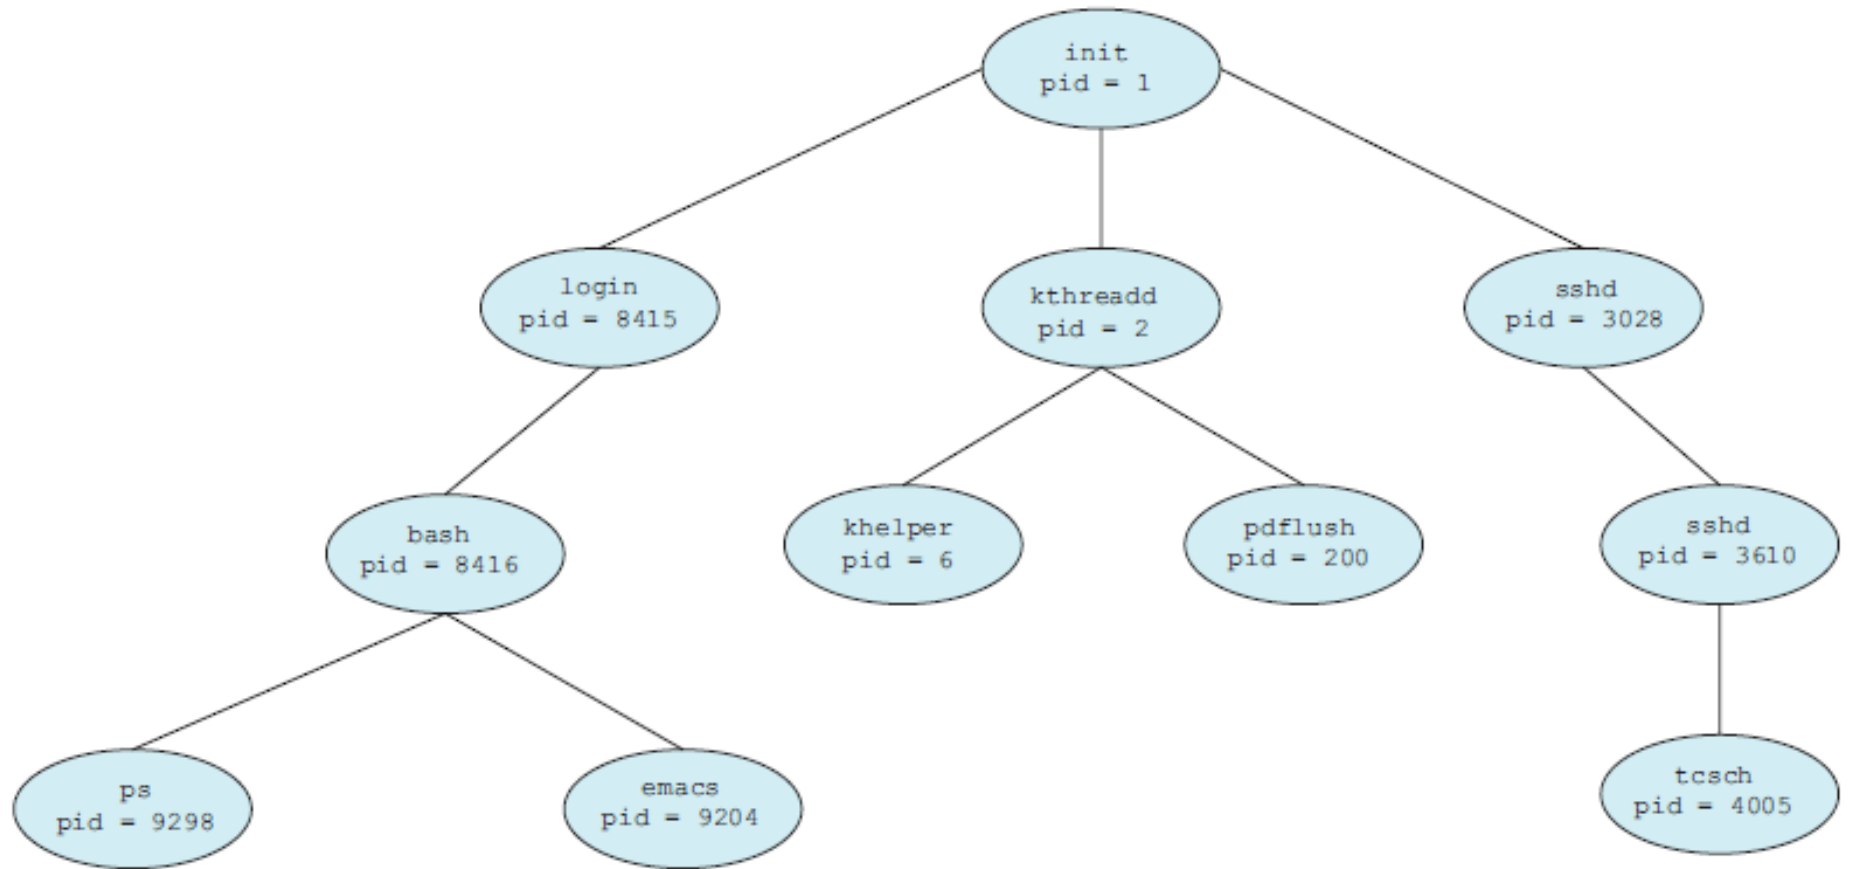
Operations on Process

- Most systems run processes dynamically and for this to be implemented, multiple processes needs to be created and destroyed.
- Broadly operations on process include:
 - Process Creation
 - Process Destruction

Process Creation

- During the course of execution a process may create multiple new processes.
- Creating process is parent process and new process created is child process.
- Operating Systems (UNIX, Linux and Windows) identify process using a unique identifier called PID.
- PID is a integer and is used as index to various attributes of a process.
-

A tree of Processes in a Linux System



Process Creation

- When a new process is created there may exist some scenarios:
 - The child program may request the required resources directly from the OS.
 - Parent may partition its resources among its children.
- Partitioning the resources avoids over burdening the OS
- Also two possibilities of execution exist:
 - The parent continues to execute concurrently with its children.
 - The parent waits until some or all of its children have terminated.
- Two address-space possibilities for the new process:
 - The child process is a duplicate of the parent process (it has the same program and data as the parent).
 - The child process has a new program loaded into it.

Process Creation

- A new process is created by a `fork()` system call (UNIX).
- New process consists of a copy of the address space of the original process, which allows communication between parent and child.
- Both parent and child execute after `fork()`
 - Child process has a return identifier 0.
 - Parent process has a non-negative return identifier
- Child process is an identical copy of parent process.
- The new processes can start execution using `exec()` system call.
- In PCB every process has a reference to its parent process, hence sharing of resources is possible.
- If a parent process has nothing to do and is waiting for child to finish then it can enter to waiting using `wait()` sys call.

fork() Demonstration

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }

    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls");
    }
    • }
    • else { /* parent process */
    • /* parent will wait for the child to complete */
    • wait();
    • printf("Child Complete");
    • }
    • return 0;
    • }
```

Process Termination

- Process terminates when it finishes executing.
- Terminates using `exit()` system call.
- A terminating process informs its parent by returning an integer.
- All the resources i.e. physical and virtual memory, open files, and I/O buffers are deallocated by OS.
- A parent can also abruptly terminate a process based on following reasons:
 - Child has exceeded its usage of some of the resources.
 - The task assigned to the child is no longer required.
 - The parent is exiting, and the operating system does not allow a child to continue if its parent terminates

Class Activity: For the following scenario draw a queuing diagram and demonstrate creation of 5 different processes after bootstrap (Assume it process P0).

Process	Time	Event/Resources	Running Time
P1	6	I/O Request	3
P2	3	System Call	2
P3	7	CPU Bound	7
P4	1	I/O Request	1

Home Work: Demonstrate process creation in a windows environment.

Process	Req. Running Time	Initial Queue	Second Queue	Actual Running Time	Enters to	Remaining running time
P1	6	Job Queue	Ready Queue	3	I/O Queue	3
P2	3	-	-	-	-	-
-	-	-	-	-	-	-
-	-	-	-	-	-	-
-	-	-	-	-	-	-